

COPENHAGEN BUSINESS ACADEMY



Concurrency and threads

Jens Egholm Pedersen
<jeep@cphbusiness.dk>

About me

- Political science from Aarhus University
- Software development from IT-University
- 4 years of work experience
 - Mainly in Java
- Generally a nice guy
 - Don't be afraid to ask questions

About these lectures

- One-way communication
 - You receive information
 - Make sure you understand it!
- Exploit what we prepared for you
 - Bloom's Taxonomy
 - Lecture = Understanding
 - Lecture + Preparation = Analyzing
 - Lecture + Preparation + Exercises = Creating
- When studying for the exam use 'see also'
 - **Not part of the curriculum!**

See also: [Something to read](#), [Bloom's taxonomy](#)

What you should know

Goal of today's and tomorrow's lectures on threads

- Understand concurrency and develop concurrent applications
- Use mechanisms to synchronise threads
- Understand and identify deadlocks
- Understand and use the publish-subscribe pattern in Java

Litterature: [Introduction to Java threads](#)

An instruction

- Basic building block of software
- Written by the programmer, translated into machine code

10 + 10 // 20

See also: [Instruction Set example](#)

A core

- Processes instructions
- One instruction at the time
 - One core: speed x 1
 - Two cores: speed x 2
 - Four cores: speed x 4
 - ... Not entirely true

See also: [Central Processing Unit \(CPU\)](#)

A thread


- A list of many instructions
- Can be assigned to a core
- Can start and stop



See also: [Java API for Thread class](#)

Threads in Java 1/2

- Can be created, started and stopped
- Creating threads
 - With **Thread**
 - With **Runnable**
- Who decides when a thread is run?



A scheduler

- Schedules when threads should be run
 - ‘God mode’ - you have almost no control
- Consequences:
 - Nondeterministic
 - Overhead
 - Scheduling is work
 - More threads, more work
 - Hint: don’t create 1.000.000 threads
 - Unless you have 1.000.000 cores

See also: [Definition of ‘scheduling’](#), [Almdahl’s law](#)

Multi-core architecture

- Having many cores in one CPU
 - One core: one thread at once
 - Two cores: two threads at once
 - 1.000.000 cores: 1.000.000 threads at once
 - One core trying to run 1.000.000 threads = disaster
- Many cores but shared hardware
- What are the threads doing?
 - Manipulating memory and interacting with I/O devices

See also: [Multicore processors](#)

Recap


- A thread is a list of instructions
 - A core can execute one thread at the time
 - No one knows *when* threads starts or stops
 - Threads share memory
-
- A process can create and use threads

Coming up

- How and why concurrency ~~can~~ will fail
- More threading in Java
- Synchronising threads
- Deadlock and starvation

Processes and threads

- A process is a running application
- A process can have one or more threads
- Every process has a main thread
 - In Java: `public static void main(String[] args)`



See also: [Definition of process](#)

A process in Java

- One thread
- Two threads
- Many threads

Concurrency

- “Happening at the same time”
- Normal program: A – Z
- Concurrent program: G, S, A, D, Y, ...
- What could go wrong?
- This is very very very hard

See also: [Concepts of Concurrent programming](#) (glossary), [Rust Concurrency](#)

Thread unpredictability

- Threads can start and stop **at all times**
 - Even between instructions!
- How many instructions are `count++`?
 - 3: Load, sum, store

See also: [java.util.concurrent.atomic](#) package in Java

Problems with concurrency

- Race condition
- Deadlocking
- Starvation

See also: [java.util.concurrent.atomic](#) package in Java

Race condition

- Problem with shared memory
 - Who comes first?
- Happens when order of threads matter
- Example: Summing numbers in Java
- Good news: Can be avoided with practice

See also: [Java Thread synchronisation tutorial](#)

Deadlocking in Java

- When two threads wait for the same thing
- Example: deadlock in Java
- Bad news: No easy solution
 - Advice from book

‘To avoid deadlock, you should ensure that when you acquire multiple locks, you always acquire the locks in the same order in all threads.’

See also: [Java deadlock tutorial](#)

Starvation in Java

- When threads cannot progress
 - Greedy threads steals resources
 - Thread slows down
- Hard to reproduce


Recap

- A thread is a list of instructions
- Threads run randomly but share memory
- Java processes can create and run threads
- Threads are dangerous!
 - They share memory and risk race conditions
 - They deadlock
 - They can steal resources from others

Coming up


- History of CPU development
- Multi-core processor architecture
- Benefits of multi threaded computing

Moore's law



See also: Moore's law


Clock frequency




See also: Technology roadmap for semiconductors

Where does the computing
power come from?

Multi-core processor



Multi-core processor




Power consumption


- High frequency = high energy consumption
- Low frequency + more cores = efficiency



Perf = 1
Power = 1



Perf = 2
Power = 4



Perf = 4
Power = 4

Parallelisation benefits

- A core can work one instruction at the time
 - 1 core = speed x 1
 - 2 cores = speed x 2
 - 4 cores = speed x 4
- ... If you parallelise your software!

Recap

- Number of cores are growing fast!
- Cores in a multi-core processor share memory
- Perfect parallelisation means a speedup factor equal to the number of cores
 - Fast!

Coming up

- Atomic variables
- Critical sections of code
- Locking and synchronisation
- More about threads
- Controlling threads in Java

Critical section

- Critical sections may not run concurrently
 - Example: variable
- Thread safety
 - A guarantee for the same behaviour with many threads
- How to control access to a section?
 - Gatekeeper / locking variable
- Who protects the protector?
 - Good question!

Locking a variable

- Atomic operations cannot be intercepted
 - Also known as thread-safety
 - Lock-free
 - Solves race-conditions for counters
- Javas atomic classes

```
AtomicInteger counter = new AtomicInteger();  
counter.get(); // 0  
counter.getAndIncrement();  
counter.get(); // 1
```

- Why 'getAndIncrement'?
 - Because increments are 3 operations

See also: [java.util.concurrent.atomic](https://docs.oracle.com/javase/7/docs/api/java/util/concurrent/atomic/package-summary.html)

Locking a section

- Why not just a simple variable?
 - Because it is not thread safe!
- `volatile` keyword
 - Ensures that reads and writes are atomic
 - Is this good enough?
- No, not always good enough
 - Good for reading (loop condition)
 - Bad for counting (race-condition)

See also: [Mutual exclusion \(mutex\)](#)

Synchronisation

- Only allow one thread to access code at once
- Method synchronisation
- Block synchronisation
 - Synchronising on `this`
 - Synchronizing on object(s)
- Which one is preferred?
 - Block synchronisation. It only locks the critical section and not the whole method

See also: [Java synchronization](#)

Synchronisation question

- Is this implementation thread safe?

```
int i = 0;

synchronized void increment() {
    i++;
}
```

- No! `i` is public!
 - Synchronize does not mean you do not have to think
 - Encapsulate, encapsulate, encapsulate

Recap

- Thread is a list of instructions
- Concurrency can go wrong!
- Atomic operations
- Locks
- Synchronisation
- Thread-safety

Thread safety question

- Is this thread-safe?

```
class Singleton {  
    private static Singleton instance = null;  
    private Singleton() { /* private constructor */ }  
    public static Singleton getInstance() {  
        if (instance == null) {  
            instance = new Singleton();  
        }  
        return instance;  
    }  
}
```

- **No!** getInstance is not synchronized

Coming up


- Locking and synchronisation
- More about threads
- Controlling threads in Java
- Semaphores

Threads in Java 1/2

- Can be created, started and stopped
- Creating threads

- With **Thread**
- With **Runnable**

I lied!




See also: [Java thread states](#)

Java thread states

Threads in Java 2/2

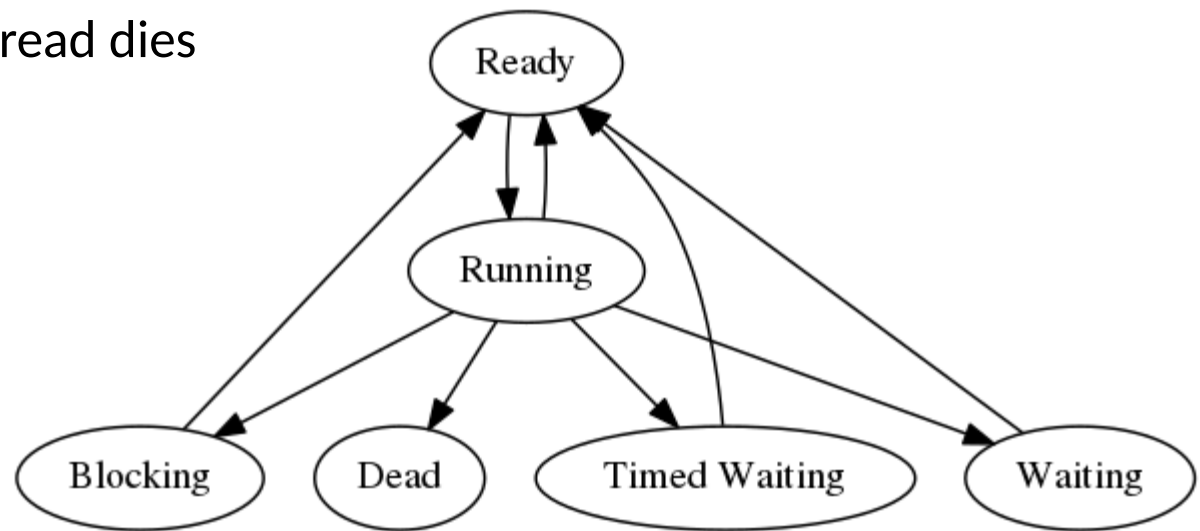
- Threads have three more states
- Blocking
 - Waiting for resource until released !
 - For instance when waiting for access to `synchronized`
- Waiting
 - `Object.wait()`;
 - `Thread.join()`;
- Waiting with a timeout
 - `Object.wait(...)`;
 - `Thread.join(...)`;
 - `Thread.sleep(...)`;



See also: [Java API for Thread class](#)

Controlling threads 1/2


- Threads can be controlled (a little)
 - `interrupt()`
 - Sends a signal to the thread to stop what it's doing
 - Can be ignored by the thread
 - `join()` and `join(...)`
 - Wait until a thread dies
 - `sleep(...)`
 - Do nothing



See also: [Java API for Thread class](#)

Controlling threads 2/2

- Threads can also be controlled via `Object`
 - `wait()` and `wait(...)`
 - Blocks current thread until someone calls `notify()`
 - `notify()`
 - Wakes a single thread
 - `notifyAll()`
 - Wakes all threads
 - Both are tied to an object
 - MUST be synchronized



See also: [Java API for Thread class](#)

A Semaphore

- Data used to control access for multiple threads
- Useful metaphor: How many units are available?
- Two types of semaphores
 - Binary (1 resource)
 - Counting (2+ resources)
- Is the `synchronized` keyword a semaphore?

See also: [Java API for Semaphore](#)

Semaphore examples

- Two operations
 - `acquire()` access to critical region
 - `release()` access permission
- Binary semaphore
 - `Semaphore s = new Semaphore(1);`
- Counting semaphore
 - `Semaphore s = new Semaphore(10);`

See also: [Java API for Semaphore](#)

Locks in Java

- Another way to controls access to critical section
- Fairness
 - Longest waiting thread gets the lock
- Lock interface
 - `lock()`
 - `unlock()`
- ReadWriteLock interface
 - `readLock()` - returns a Lock
 - `writeLock()` - returns a Lock
- Remember to release the lock!
 - Use `try-finally`

See also: [Java API for Lock](#), [Java API for ReadWriteLock](#)

Locks in Java example

- ReentrantLock

- Reentrant means that a thread can have more than one lock

```
ReentrantLock lock = new ReentrantLock();  
lock.lock();  
  
try { ... } finally { lock.unlock(); }
```

- ReentrantReadWriteLock

```
ReentrantReadWriteLock lock = new  
ReentrantReadWriteLock();  
  
WriteLock writeLock = lock.writeLock();  
writeLock.lock();  
  
try { ... } finally { writeLock.unlock(); }
```

See also: [ReentrantReadWriteLock versus StampedLock](#)

Synchronisation versus locks

- Lock provides more visibility and functionality
 - For instance `tryLock()`
- Lock **require** `try-finally` blocks
 - Synchronization code can be cleaner
- Lock can be aquired in one method and released in another
 - Good or bad?
- Lock provide fairness
 - The longest waiting threads get the lock first

Recap

- Critical sections consist of code that cannot run concurrently
- Atomic operations cannot be intercepted
- Controls to critical sections
 - `synchronized`
 - Semaphores (can allow one or more threads at once)
 - Lock
- Threads can be `join()`'ed and `interrupt()`'ed

Coming up

- Thread priority
- Thread pools and `ExecutorService` in Java
- Threads in Swing
- Producer-consumer problem

Thread priority

- Thread can have a priority
 - Tells the scheduler how to prioritise
- Ranges from 1 (min) to 10 (max)
 - `Thread.MIN_PRIORITY` and `Thread.MAX_PRIORITY`
- Must be set *before* start

```
Thread t = ...;  
t.setPriority();  
t.start();
```

Executors in Java

- Creating threads is not free
 - Reuse them!
- Thread pool in Java

Executors


```
.newCachedThreadPool()  
.execute(runnable);
```

Threads in Swing

- Swing uses a special event dispatch thread
- All Swing component methods must be invoked from this thread
 - Unless the API states it is thread safe
- Example
 - `SwingUtils.invokeLater(() -> text.setText("hi"));`
- What if you submit a 1000 second task?
 - It blocks!

Producer-consumer problem 1/2

- Consider two processes:
 - One process puts elements in a queue
 - One process takes elements from the queue
- The queue has a fixed size
- What happens if
 - The producer slows down?
 - Starvation
 - The consumer slows down?
 - Buffer overflow



Producer-consumer problem 2/2

- In modern Java: `ArrayBlockingQueue`
- Example

```
ArrayBlockingQueue q = new ArrayBlockingQueue();  
// Thread 1 - Producer  
q.put(...);  
// Thread 2 - Consumer  
q.take();
```

Detecting deadlocks

```
ThreadMXBean mx =  
    ManagementFactory.getThreadMXBean();  
long[] deadlocked = mx.findDeadlockedThreads();
```


Recap

- Threads can be prioritised
- Use executors when you have many tasks
- Swing uses one single event thread
- Producer-consumer problem

Exercises

- Exercises on Fronter